

Unit 11. Decision Trees

1 Cost-Complexity Pruning

ISLR §8.1.1, tree pruning (both reading lists). SRM learning outcome 4(a).

Remark (ISLR §8.1.1). A tree grown until every leaf is nearly pure fits the training data and overfits. A threshold on the split gain is not the repair: a split worth little on its own can make the next split worth a great deal, and a threshold discards it first. ISLR's remedy, Algorithm 8.1, is to grow a large tree and prune it.

1.1 The pruning family

Definition 11.1 (Tree, branch, pruned subtree). $T_{\max}$ is a finite binary tree: each node t has either two children t_L, t_R or none, in which case it is a **leaf** of $T_{\max}$. The **branch** $T_{\max, t}$ is the set of descendants of t , including t ; ISLR calls the segments joining the nodes branches, and here, as in the pruning literature, the branch at t in §1 is the whole subtree below it. A **pruned subtree rooted at** t , written $S \preceq T_{\max, t}$, is obtained from $T_{\max, t}$ by choosing a set of nodes and deleting everything strictly below each of them; $\mathcal{S}(t)$ denotes the collection of these, a finite set containing both $\{t\}$ and $T_{\max, t}$. More generally, for $S', S \in \mathcal{S}(t)$, write $S' \preceq S$ if every node of S' is a node

of S (equivalently, S' is obtained from S in the same way), and $S' \prec S$, or $S \succ S'$, if moreover $S' \neq S$. Write $|S|$ for the number of leaves of S .

Definition 11.2 (Node cost). A **node cost** assigns $R(t) \geq 0$ to every node, the cost incurred if t is made a leaf, and a tree is charged the sum over its leaves,

$$R(S) = \sum_{t \text{ leaf of } S} R(t). \quad (1)$$

It satisfies the **refinement property** if $R(t) \geq R(t_L) + R(t_R)$ at every internal node.

Proposition 11.3 (The three costs of this unit refine). Let node t hold $n_t \geq 1$ observations, and each of its children at least one. Each of

- (i) the residual sum of squares $R(t) = \sum_{i \in t} (y_i - \bar{y}_t)^2$,
- (ii) the misclassification count $R(t) = n_t - \max_k n_{tk}$, where n_{tk} is the number of class k observations in t ,
- (iii) $R(t) = n_t \phi(\mathbf{p}_t)$ for any concave $\phi \geq 0$ on the probability simplex, $\mathbf{p}_t$ being the vector of

class proportions in t ,
satisfies the refinement property, and (ii) is the case $\phi(\mathbf{p}) = 1 - \max_k p_k$ of (iii).

Proof. (i) is the analysis of variance identity of Unit 10, whose third term $n_L n_R (\bar{y}_L - \bar{y}_R)^2 / n_t$ is non-negative.

(iii) The class counts add, $n_{tk} = n_{Lk} + n_{Rk}$, so $\mathbf{p}_t = (n_L/n_t)\mathbf{p}_L + (n_R/n_t)\mathbf{p}_R$, and concavity gives $\phi(\mathbf{p}_t) \geq (n_L/n_t)\phi(\mathbf{p}_L) + (n_R/n_t)\phi(\mathbf{p}_R)$; multiplying by n_t gives $R(t) \geq R(t_L) + R(t_R)$.

(ii) $\mathbf{p} \mapsto \max_k p_k$ is a maximum of linear functions, hence convex, so $1 - \max_k p_k$ is concave and (iii) applies. Directly: $\max_k (n_{Lk} + n_{Rk}) \leq \max_k n_{Lk} + \max_k n_{Rk}$. $\square$

Definition 11.4 (Cost complexity; ISLR (8.4)). For $\alpha \geq 0$ and $S \in \mathcal{S}(t)$,

$$R_\alpha(S) = R(S) + \alpha |S|. \quad (2)$$

Remark. The number of pruned subtrees can grow faster than any polynomial in the size of $T_{\max}$: a complete binary tree of depth d has N_d pruned subtrees with $N_0 = 1$ and $N_d = 1 + N_{d-1}^2$, since such a subtree is either the root alone or a pair of pruned subtrees of the two branches. That gives

$N_5 = 458,330$ and $N_6 > 2 \times 10^{11}$ on a tree of 127 nodes. Minimising (2) by enumeration is not available.

1.2 The minimising subtree

Definition 11.5 (Value function). $m_\alpha(t) = \min_{S \in \mathcal{S}(t)} R_\alpha(S)$.

Lemma 11.6 (Recursion). If t is a leaf of $T_{\max}$ then $m_\alpha(t) = R(t) + \alpha$. Otherwise

$$m_\alpha(t) = \min \{ R(t) + \alpha, \quad m_\alpha(t_L) + m_\alpha(t_R) \}. \quad (3)$$

Proof. A pruned subtree rooted at an internal t either is $\{t\}$, of cost complexity $R(t) + \alpha$, or retains both children, in which case it is determined by a pair $S_L \in \mathcal{S}(t_L)$, $S_R \in \mathcal{S}(t_R)$ and its leaves are those of S_L together with those of S_R . Both R and $|\cdot|$ add over that partition of the leaves, so $R_\alpha(S) = R_\alpha(S_L) + R_\alpha(S_R)$, and minimising over the pair minimises the two terms separately. $\square$

Lemma 11.7 (Shape in α). $\alpha \mapsto m_\alpha(t)$ is the minimum of the finitely many affine functions $\alpha \mapsto R(S) + \alpha|S|$, $S \in \mathcal{S}(t)$. It is therefore continuous, piecewise linear, concave, and strictly increasing, every slope $|S|$ being at least 1.

Proof. Definition 11.5 is exactly that minimum, over a finite set. A minimum of finitely many affine functions is concave and piecewise linear, and it is strictly increasing because each of them is. $\square$

Theorem 11.8 (Every internal node has a pruning threshold). Assume the refinement property and let t be internal in $T_{\max}$. Put

$$h_t(\alpha) = m_\alpha(t_L) + m_\alpha(t_R) - R(t) - \alpha. \quad (4)$$

Then h_t is continuous, piecewise linear and strictly increasing with every slope at least 1, $h_t(0) \leq 0$, and $h_t(\alpha) \rightarrow \infty$. Hence there is a unique $\alpha_t \in [0, \infty)$ with

$$h_t(\alpha) < 0 \text{ for } \alpha < \alpha_t, \quad h_t(\alpha) \geq 0 \text{ for } \alpha \geq \alpha_t. \quad (5)$$

Proof. By Lemma 11.7 each of $m_\alpha(t_L)$ and $m_\alpha(t_R)$ is piecewise linear with slopes at least 1, so their sum has slopes at least 2 and h_t has slopes at least $2 - 1 = 1$: it is continuous, piecewise linear, strictly increasing, and unbounded above.

At $\alpha = 0$ the minimum in Definition 11.5 is over $R(S)$ alone, and iterating the refinement property down the branch shows $R(T_{\max,t}) \leq R(S)$ for every $S \in \mathcal{S}(t)$, so $m_0(t) = R(T_{\max,t})$. Hence $h_t(0) = R(T_{\max,t_L}) + R(T_{\max,t_R}) - R(t) = R(T_{\max,t}) - R(t) \leq 0$, again by refinement. A continuous strictly increasing function that is non-positive at 0 and unbounded above has a unique zero $\alpha_t \geq 0$, and (5) follows. $\square$

Definition 11.9 (The canonical subtree). Assume the refinement property, so that every α_t of Theorem 11.8 exists; it is assumed to the end of §1. $S_\alpha(t)$ is defined upward from the leaves: $S_\alpha(t) = \{t\}$ if t is a leaf of $T_{\max}$ or if $\alpha \geq \alpha_t$; otherwise $S_\alpha(t)$ is t together with $S_\alpha(t_L)$ and $S_\alpha(t_R)$. Write $T(\alpha) = S_\alpha(r)$ at the root r .

Theorem 11.10 ($T(\alpha)$ is the smallest minimiser). For every $\alpha \geq 0$ and every node t ,

$$R_\alpha(S_\alpha(t)) = m_\alpha(t), \quad \text{and} \quad S_\alpha(t) \preceq S \text{ for every minimiser } S \in \mathcal{S}(t). \quad (6)$$

In particular $T(\alpha)$ minimises R_α over $\mathcal{S}(r)$ and is a pruned subtree of every other minimiser.

Proof. Induction on the height of t . If t is a leaf of $T_{\max}$ then $\mathcal{S}(t) = \{\{t\}\}$ and both statements are trivial.

Let t be internal. If $\alpha \geq \alpha_t$ then $h_t(\alpha) \geq 0$, that is $R(t) + \alpha \leq m_\alpha(t_L) + m_\alpha(t_R)$, so (3) gives $m_\alpha(t) = R(t) + \alpha = R_\alpha(\{t\})$; and $S_\alpha(t) = \{t\}$ is a pruned subtree of every member of $\mathcal{S}(t)$.

If $\alpha < \alpha_t$ then $h_t(\alpha) < 0$, so $m_\alpha(t) = m_\alpha(t_L) + m_\alpha(t_R)$ and the inequality $R(t) + \alpha > m_\alpha(t)$ is strict: no minimiser is $\{t\}$. Every minimiser S therefore retains both children and splits as a pair (S_L, S_R) with $R_\alpha(S_L) + R_\alpha(S_R) = m_\alpha(t_L) + m_\alpha(t_R)$, which forces S_L and S_R to be minimisers at t_L and t_R . By the induction hypothesis $R_\alpha(S_\alpha(t_L)) + R_\alpha(S_\alpha(t_R)) = m_\alpha(t)$, which is the first claim, and $S_\alpha(t_L) \preceq S_L$, $S_\alpha(t_R) \preceq S_R$, which is the second. $\square$

Theorem 11.11 (Nesting). If $0 \leq \alpha < \alpha'$ then $T(\alpha') \preceq T(\alpha)$.

Proof. Induction on the height of t , proving $S_{\alpha'}(t) \preceq S_\alpha(t)$. For a leaf of $T_{\max}$ both are $\{t\}$. For internal t , if $\alpha' \geq \alpha_t$ then $S_{\alpha'}(t) = \{t\}$, which is a pruned subtree of anything. Otherwise $\alpha < \alpha' < \alpha_t$, so both $S_{\alpha'}(t)$ and $S_\alpha(t)$ retain both children, and the induction hypothesis applies

at t_L and t_R . □

Corollary 11.12 (A finite nested sequence; ISLR §8.1.1). There are α values $0 = \eta_0 < \eta_1 < \dots < \eta_B$ and trees $T(0) = T^{(0)} \succ T^{(1)} \succ \dots \succ T^{(B)} = \{r\}$ with $T(\alpha) = T^{(j)}$ for $\alpha \in [\eta_j, \eta_{j+1})$ and $T(\alpha) = T^{(B)}$ for $\alpha \geq \eta_B$, where $T(0)$, the smallest minimiser of R , is $T_{\max}$ when $R(T_{\max,t}) < R(t)$ at every internal t . Moreover $B < |T_{\max}|$, so cross-validation over α is a search over at most $|T_{\max}|$ candidate trees.

Proof. By Definition 11.9, $T(\alpha)$ depends on α only through which of the finitely many inequalities $\alpha \geq \alpha_t$ hold, so it is piecewise constant with breakpoints among $\{\alpha_t\}$. By Theorem 11.11 the distinct values are strictly nested, and each strict pruning removes at least one leaf, so there are at most $|T_{\max}|$ of them. At $\alpha = 0$ the definition expands exactly where $\alpha_t > 0$, that is where $R(T_{\max,t}) < R(t)$ (Theorem 11.8: $\alpha_t = 0$ exactly when $h_t(0) = 0$, and its proof gives $h_t(0) = R(T_{\max,t}) - R(t)$), so the value is $T_{\max}$ when this holds at every internal t ; for α past every α_t it is $\{r\}$. □

Corollary 11.13 (Weakest link; ISLR §8.1.1). Let t be internal and let T_t be t together with $S_{\alpha_t}(t_L)$ and $S_{\alpha_t}(t_R)$; it has $|T_t| \geq 2$ leaves, at least one from each side. Then

$$\alpha_t = \frac{R(t) - R(T_t)}{|T_t| - 1}. \quad (7)$$

Proof. $h_t(\alpha_t) = 0$ reads $m_{\alpha_t}(t_L) + m_{\alpha_t}(t_R) = R(t) + \alpha_t$, and by Theorem 11.10 the left side is $R_{\alpha_t}(T_t) = R(T_t) + \alpha_t|T_t|$. Solve for α_t . $\square$

Remark. Equation (7) is the quantity every implementation reports: the cost saved by keeping the branch, per extra leaf. `tree` prints it as the vector `k` of a pruning sequence, which lists the α_t at which $T(\alpha)$ changes; `rpart` prints `cp`, which is (7) divided by the cost of the root, so that `cp` is a fraction of the total.

Remark. The proof is an algorithm. One sweep from the leaves upward evaluates m_α and $T(\alpha)$ at a given α by (3) in a number of operations proportional to the number of nodes, against a count of pruned subtrees that can grow exponentially in the number of nodes.

1.3 Choosing the penalty

Definition 11.14 (Cross-validation over the sequence; ISLR Algorithm 8.1). Split the data into K folds. For each fold k , grow a tree on the other $K - 1$ folds, form its own pruning sequence, and record the cost on fold k of the subtree selected at each α in a common grid. Average over folds to obtain $CV(\alpha)$ and take $\hat{\alpha} = \arg \min_{\alpha} CV(\alpha)$. The tree reported is $T(\hat{\alpha})$ from the sequence grown on all the data.

Remark. The grid must be common to all folds, and the fold sequences have their own break-points, so the grid is taken from the full-data sequence. Two consequences are easy to overlook. The subtree a fold contributes at a given α need not have the number of leaves that the full-data tree has at that α , so the size axis of a cross-validation plot is a label from the full data and not a property of what was scored; and the whole tree, not only its pruning, must be regrown inside each fold, since choosing the splits on all the data and cross-validating only the pruning would evaluate a rule that has already seen the held-out observations.

Definition 11.15 (One standard error rule; Unit 5). Let $\widehat{\text{se}}(\alpha)$ be the standard error over folds of the cross-validated cost. The **one standard error rule** selects the largest α , that is the smallest tree, with

$$\text{CV}(\alpha) \leq \text{CV}(\hat{\alpha}) + \widehat{\text{se}}(\hat{\alpha}). \quad (8)$$

Remark. Corollary 11.12 makes Definition 11.14 feasible: the candidate set is a nested sequence of at most $|T_{\max}|$ trees, not the collection of all pruned subtrees. Nothing here shows that $\hat{\alpha}$ estimates the minimiser of the test error well; that is the cross-validation question of Unit 1.

2 Classification Trees

ISLR §8.1.2. SRM learning outcomes 4(a), 4(b).

2.1 Leaf labels

Definition 11.16 (Classification tree; ISLR §8.1.2). The response takes values in $\{1, \dots, K\}$; from here on K counts classes, not folds. A classification tree assigns to each leaf t the class proportions $\mathbf{p}_t = (p_{t1}, \dots, p_{tK})$ with $p_{tk} = n_{tk}/n_t$, and predicts a single class $\hat{k}(t)$ for every observation falling in t .

Proposition 11.17 (Majority vote; ISLR §8.1.2). The label minimising the number of training errors in leaf t is any modal class $\hat{k}(t) \in \arg \max_k n_{tk}$, and the resulting error count is $n_t - \max_k n_{tk} = n_t M(\mathbf{p}_t)$ with $M(\mathbf{p}) = 1 - \max_k p_k$.

Proof. Labelling t with class k misclassifies the $n_t - n_{tk}$ observations of other classes, which is smallest when n_{tk} is largest. $\square$

Remark. Unit 1 proved that the classifier minimising the error probability assigns $\mathbf{x}$ to $\arg \max_k \mathbf{P}\{Y = k \mid X = \mathbf{x}\}$. Proposition 11.17 is that rule with the conditional probabilities replaced by the proportions in the leaf, so a classification tree is an estimate of the Bayes classifier in which the conditioning event is the box rather than the point. The class proportions

themselves, not only the label, are part of the output: two leaves may share a label and differ entirely in how confident it is.

2.2 Impurity

Definition 11.18 (Impurity function). An **impurity** is a continuous concave $\phi : \Delta^{K-1} \rightarrow [0, \infty)$ on the probability simplex $\Delta^{K-1} = \{\mathbf{p} \geq \mathbf{0} : \sum_k p_k = 1\}$ with $\phi(\mathbf{e}_k) = 0$ at every vertex. The impurity of a tree is $R(T) = \sum_{t \text{ leaf}} n_t \phi(\mathbf{p}_t)$ and the **gain** of a split is

$$\Delta_\phi = \phi(\mathbf{p}_t) - \frac{n_L}{n_t} \phi(\mathbf{p}_L) - \frac{n_R}{n_t} \phi(\mathbf{p}_R). \quad (9)$$

Definition 11.19 (The three impurities; ISLR (8.5)–(8.7)).

$$M(\mathbf{p}) = 1 - \max_k p_k, \quad G(\mathbf{p}) = \sum_k p_k(1 - p_k) = 1 - \sum_k p_k^2, \quad H(\mathbf{p}) = - \sum_k p_k \ln p_k, \quad (10)$$

the **misclassification rate**, the **Gini index** and the **entropy**, with $0 \ln 0 = 0$.

Proposition 11.20 (Concavity). M , G and H are impurities. G and H are strictly concave on the simplex; M is concave but affine on every region where a single class is strictly modal.

Proof. All three vanish at the vertices and are continuous, $x \ln x \rightarrow 0$ as $x \downarrow 0$ giving continuity of H .

G : along every line, $u \mapsto G(\mathbf{p} + u\mathbf{v})$ has second derivative $-2\|\mathbf{v}\|^2 < 0$ for $\mathbf{v} \neq \mathbf{0}$, so G is strictly concave on the simplex by the lemma on convexity from the derivative (prerequisites), which reduces strict concavity to restrictions to lines.

H : $\psi(x) = -x \ln x$ is continuous on $[0, 1]$ with $\psi'(x) = -\ln x - 1$ strictly decreasing on $(0, 1)$, so ψ is strictly concave on the closed interval $[0, 1]$ by the same lemma. If $\mathbf{p} \neq \mathbf{q}$ in the simplex, some coordinate differs, so for $0 < \lambda < 1$ every term satisfies $\psi(\lambda p_k + (1 - \lambda)q_k) \geq \lambda \psi(p_k) + (1 - \lambda)\psi(q_k)$, strictly for that coordinate; summing gives $H(\lambda \mathbf{p} + (1 - \lambda)\mathbf{q}) > \lambda H(\mathbf{p}) + (1 - \lambda)H(\mathbf{q})$, boundary points included.

M : $\max_k p_k$ is a maximum of linear functions, hence convex, so M is concave. On the set where class k^* is strictly modal, $M(\mathbf{p}) = 1 - p_{k^*}$ is affine. $\square$

Theorem 11.21 (A split never hurts, and when it does nothing). For any impurity ϕ , $\Delta_\phi \geq 0$. If ϕ is strictly concave and both children are non-empty, $\Delta_\phi = 0$ if and only if $\mathbf{p}_L = \mathbf{p}_R$.

Proof. The class counts add, so $\mathbf{p}_t = (n_L/n_t)\mathbf{p}_L + (n_R/n_t)\mathbf{p}_R$ and concavity gives $\phi(\mathbf{p}_t) \geq (n_L/n_t)\phi(\mathbf{p}_L) + (n_R/n_t)\phi(\mathbf{p}_R)$, which is (9) ≥ 0 . Strict concavity makes the inequality strict unless $\mathbf{p}_L = \mathbf{p}_R$, in which case both sides equal $\phi(\mathbf{p}_t)$. $\square$

Theorem 11.22 (Misclassification is blind to a split that keeps the modal class). If the same class k^* is modal in the parent, in the left child and in the right child, then $\Delta_M = 0$ exactly.

Proof. Under the hypothesis, $n_t M(\mathbf{p}_t) = n_t - n_{tk^*}$ and likewise for the children, and the counts add: $n_L M(\mathbf{p}_L) + n_R M(\mathbf{p}_R) = (n_L - n_{Lk^*}) + (n_R - n_{Rk^*}) = n_t - n_{tk^*} = n_t M(\mathbf{p}_t)$. Divide by n_t . $\square$

Remark. Trees are therefore not grown on the error count. A split that leaves both children with the same modal class scores exactly zero on M however differently the two children are composed, so a greedy search on M has no gradient to follow and stops. Under G or H the same split has

strictly positive gain unless the two child distributions coincide (Theorem 11.21), and the search continues.

Theorem 11.23 (Ordering). For every $\mathbf{p} \in \Delta^{K-1}$,

$$M(\mathbf{p}) \leq G(\mathbf{p}) \leq H(\mathbf{p}), \quad (11)$$

and for $K = 2$ also $G(\mathbf{p}) \leq 2M(\mathbf{p})$, with equality only at the vertices (while $M = G$ at $\mathbf{p} = (\frac{1}{2}, \frac{1}{2})$).

Proof. For the first inequality, $\sum_k p_k^2 \leq (\max_k p_k) \sum_k p_k = \max_k p_k$, so $G = 1 - \sum_k p_k^2 \geq 1 - \max_k p_k = M$.

For the second, $\ln u \leq u - 1$ for $u > 0$, applied at $u = p_k$, gives $-\ln p_k \geq 1 - p_k$ and hence $H = \sum_k p_k(-\ln p_k) \geq \sum_k p_k(1 - p_k) = G$, the terms with $p_k = 0$ contributing 0 to both sides.

For $K = 2$ write $\mathbf{p} = (p, 1-p)$ and $q = \min\{p, 1-p\}$, so $M = q$ and $G = 2p(1-p) = 2q(1-q) \leq 2q$ since $1-q \leq 1$. Equality needs $q^2 = 0$, that is $q = 0$, where both sides vanish. $\square$

Proposition 11.24 (What the Gini index measures). Let U, V be independent draws from $\mathbf{p}$. Then $G(\mathbf{p}) = \mathbf{P}\{U \neq V\}$: it is the error rate of the rule that guesses a class at random from the leaf's own proportions. For $K = 2$, $G(\mathbf{p}) = 2 \operatorname{var}(\mathbb{1}\{U = 1\})$.

Proof. $\mathbf{P}\{U = V\} = \sum_k p_k^2$. For $K = 2$, $\operatorname{var}(\mathbb{1}\{U = 1\}) = p(1 - p)$ and $G = 2p(1 - p)$. $\square$

Proposition 11.25 (Entropy and the reported deviance; ISLR §8.3.1).

The quantity $-2 \sum_{t \text{ leaf}} \sum_k n_{tk} \ln p_{tk}$, which `tree` prints as the deviance of a classification tree, equals $2 \sum_{t \text{ leaf}} n_t H(\mathbf{p}_t)$. It is the deviance of Unit 8 (the definition of deviance there): twice the log likelihood ratio of the saturated model, which gives each observation its own class probability one, against the fitted leaf probabilities.

Proof. $\sum_k n_{tk} \ln p_{tk} = n_t \sum_k p_{tk} \ln p_{tk} = -n_t H(\mathbf{p}_t)$. For the second sentence, the log likelihood of the individual observations in leaf t is $\sum_k n_{tk} \ln p_{tk}$, and the saturated model puts each observation's own class at probability one, with log likelihood 0. $\square$

Remark. The three quantities therefore serve different roles. G or H grows the tree, because Theorem 11.22 says M cannot. M prunes it, because Corollary 11.12 needs a cost and the cost one actually cares about is the error count. Both are legitimate node costs by Proposition 11.3, and the two choices can select trees of very different sizes on the same data.

2.3 Two classes

Theorem 11.26 (A two-class Gini split is a regression split on the indicator). Let $K = 2$ and code $y_i = \mathbf{1}\{\text{observation } i \text{ is in class 2}\}$. Then for any split of node t ,

$$n_L G(\mathbf{p}_L) + n_R G(\mathbf{p}_R) = 2 \left\{ \sum_{i \in L} (y_i - \bar{y}_L)^2 + \sum_{i \in R} (y_i - \bar{y}_R)^2 \right\}, \quad (12)$$

so minimising the weighted Gini impurity and minimising the residual sum of squares of the 0–1 response select the same split, and the Gini gain (9) is $2n_L n_R (\bar{y}_L - \bar{y}_R)^2 / n_t^2$, which is $2/n_t$ times Unit 10's gain.

Proof. In a node holding m observations, with a proportion $\bar{y}$ of them coded 1, $\sum_i (y_i - \bar{y})^2 = m\bar{y}(1 - \bar{y})$ because $y_i^2 = y_i$. The class proportions are $\mathbf{p} = (1 - \bar{y}, \bar{y})$, so $G(\mathbf{p}) = 2\bar{y}(1 - \bar{y})$ and $mG(\mathbf{p}) = 2\sum_i (y_i - \bar{y})^2$. Applying this in each child gives (12). The gain expression is the analysis of variance identity of Unit 10 applied to the 0–1 response, multiplied by $2/n_t$. $\square$

Remark. Nothing new has to be implemented for two-class Gini splitting: it is the regression splitter of Unit 10 run on the indicator. The identity also explains why the Gini gain carries Unit 10's $n_L n_R / n_t$ weight, divided by the constant $n_t/2$, so a split peeling off a handful of observations is penalised for the same reason.

2.4 Categorical predictors

Remark (Qualitative predictors; ISLR §8.1.2). A split on an unordered predictor with L levels sends a non-empty proper subset of the levels to the left, and a subset and its complement give the same split, so there are $2^{L-1} - 1$ splits: $2^{31} - 1$ at $L = 32$.

Theorem 11.27 (Ordered search suffices). Let x be categorical with levels $1, \dots, L$, let $\bar{y}_l$ be the mean response of the observations at level l , and suppose the $\bar{y}_l$ are distinct and ordered $\bar{y}_1 < \dots < \bar{y}_L$. Then some subset A^* minimising

$$Q(A) = \sum_{i: x_i \in A} (y_i - \bar{y}_A)^2 + \sum_{i: x_i \notin A} (y_i - \bar{y}_{A^c})^2 \quad (13)$$

over non-trivial A is an initial segment $\{1, \dots, m\}$. By Theorem 11.26 the same holds for a two-class Gini split, with $\bar{y}_l$ the proportion of class 2 at level l .

Proof. Write $Q(A) = \min_{a,b} F(A, a, b)$ with $F(A, a, b) = \sum_{x_i \in A} (y_i - a)^2 + \sum_{x_i \notin A} (y_i - b)^2$, the minimum over a and b being attained at the two group means. So a minimiser of Q is part of a minimiser of F over all three arguments.

Fix a and b and minimise F over A . All observations at a level move together, and level l contributes

$\sum_{i \in l} (y_i - a)^2$ if $l \in A$ and $\sum_{i \in l} (y_i - b)^2$ otherwise. Splitting each at $\bar{y}_l$,

$$\sum_{i \in l} (y_i - a)^2 - \sum_{i \in l} (y_i - b)^2 = n_l \{ (\bar{y}_l - a)^2 - (\bar{y}_l - b)^2 \} = n_l (b - a) (2\bar{y}_l - a - b),$$

the within-level sums of squares cancelling. Hence, when $a < b$, the contribution is smaller by putting l in A exactly when $\bar{y}_l < (a + b)/2$, and larger exactly when $\bar{y}_l > (a + b)/2$.

Let A^* minimise Q and put $a = \bar{y}_{A^*}$, $b = \bar{y}_{A^{*c}}$. If $a = b$ the analysis of variance identity makes $Q(A^*)$ equal to the total sum of squares in the node, so no split reduces it and every A , the initial segments included, is a minimiser. Otherwise assume $a < b$, exchanging the names of A^* and A^{*c} if necessary. Since (A^*, a, b) minimises F , no level can be moved to reduce it (a move that empties A^* or its complement gives the trivial split, whose F is at least the total sum of squares, hence at least $Q(A^*)$), so every $l \in A^*$ has $\bar{y}_l \leq (a + b)/2$ and every $l \notin A^*$ has $\bar{y}_l \geq (a + b)/2$. Since the $\bar{y}_l$ are distinct, every level in A^* therefore precedes every level outside it in the ordering: A^* is an initial segment. $\square$

Remark. The proof is the two-mean assignment step: put each level with whichever of the two fitted values it is closer to. Sorting the levels once and scanning the $L - 1$ cutpoints therefore finds an optimal split at the cost of Unit 10's continuous case. Theorem 11.27 needs a scalar $\bar{y}_l$ to sort by, so it covers regression and two-class classification; with $K > 2$ classes there is no such ordering and implementations fall back on heuristics.

3 Trees Against Linear Models

ISLR §8.1.3–8.1.4. SRM learning outcome 4(a).

Remark (ISLR (8.8)–(8.9)). A linear model fits $f(\mathbf{x}) = \beta_0 + \sum_j \beta_j x_j$ and a tree fits $f(\mathbf{x}) = \sum_m c_m \mathbf{1}\{\mathbf{x} \in R_m\}$; neither class contains the other, and which predicts better depends on the truth, as the theorems of this section quantify.

3.1 The cost of a staircase

Theorem 11.28 (Approximating a linear function by boxes). Let X be uniform on $[0, 1]$ and $f(x) = \beta_0 + \beta_1 x$. Among all partitions of $[0, 1]$ into J intervals, the mean squared error of the best piecewise-constant approximation,

$$\min_{\text{partitions}} \mathbb{E}[\{f(X) - \mathbb{E}[f(X) \mid X \in I(X)]\}^2] = \frac{\beta_1^2}{12J^2}, \quad (14)$$

is attained by the J intervals of equal width. In p dimensions, with X uniform on $[0, 1]^p$, $f(\mathbf{x}) = \beta_0 + \boldsymbol{\beta}'\mathbf{x}$ and a grid of $J = m^p$ equal boxes,

$$\mathbb{E}[\{f(X) - \mathbb{E}[f(X) \mid \text{box}]\}^2] = \frac{\|\boldsymbol{\beta}\|^2}{12 J^{2/p}}. \quad (15)$$

Proof. On an interval I of width w the conditional law of X is uniform, so $\text{var}(f(X) \mid X \in I) = \beta_1^2 w^2 / 12$, and the error contributed by I is $\mathbb{P}\{X \in I\} \text{var}(f(X) \mid X \in I) = \beta_1^2 w^3 / 12$. Summing, the error is $(\beta_1^2 / 12) \sum_j w_j^3$ subject to $\sum_j w_j = 1$. The function $u \mapsto u^3$ is convex, so Jensen's

inequality (prerequisites) gives $J^{-1} \sum_j w_j^3 \geq (J^{-1} \sum_j w_j)^3 = J^{-3}$, that is $\sum_j w_j^3 \geq J^{-2}$, with equality when all w_j are equal. That is (14).

In p dimensions the coordinates of X are independent, and on a box of side $1/m$ in every coordinate the conditional variance of $\beta'X$ is $\sum_j \beta_j^2/(12m^2)$, free of which box. With $J = m^p$, $m^2 = J^{2/p}$. $\square$

Remark. Halving the error of a linear truth costs $\sqrt{2}$ times as many leaves in one dimension. On the grid of equal boxes in p dimensions it costs $2^{p/2}$ times as many leaves: at $p = 10$ a factor of 32. When only some coordinates of β are non-zero, a tree can cut along those alone, and p is then their number. A linear model has no such error at all. The difficulty is dimensional, not a matter of sample size.

3.2 Where a tree is exactly right

Theorem 11.29 (A single step). Let X be uniform on $[0, 1]$ and $f(x) = \mathbf{1}\{x > c\}$ with $c \in (0, 1)$. The best linear predictor of $f(X)$ has mean squared error

$$c(1 - c)\{1 - 3c(1 - c)\}, \quad (16)$$

which is $1/16$ at $c = \frac{1}{2}$, while the two-leaf tree that splits at c has mean squared error 0. Moreover the split gain of Unit 10, in population form $s \mapsto \mathbf{P}\{X \leq s\} \mathbf{P}\{X > s\} \{\mathbf{E}[f \mid X \leq s] - \mathbf{E}[f \mid X > s]\}^2$, is maximised at $s = c$, so greedy splitting finds the right cut.

Proof. The best linear predictor $a + bX$ minimises $\mathbf{E}[(f - a - bX)^2] = \text{var}(f - bX) + (\mathbf{E} f - a - b \mathbf{E} X)^2$; the second term vanishes at $a = \mathbf{E} f - b \mathbf{E} X$, and $\text{var}(f) - 2b \text{cov}(f, X) + b^2 \text{var}(X)$ is least at $b = \text{cov}(f, X) / \text{var}(X)$, with value $\text{var}(f) \{1 - \text{corr}^2\}$. $\mathbf{E}[f] = 1 - c$ and $f^2 = f$, so $\text{var}(f) = c(1 - c)$; $\text{var}(X) = 1/12$ and $\text{cov}(f, X) = \int_c^1 x \, dx - (1 - c)/2 = (1 - c^2)/2 - (1 - c)/2 = c(1 - c)/2$. The squared correlation is therefore $\{c(1 - c)/2\}^2 / \{c(1 - c)/12\} = 3c(1 - c)$, and the best linear predictor leaves $\text{var}(f) \{1 - \text{corr}^2\}$, which is (16); at $c = \frac{1}{2}$ it is $\frac{1}{4}(1 - \frac{3}{4}) = \frac{1}{16}$. The tree that splits at c predicts 0 below and 1 above, reproducing f .

For the gain, take $s \leq c$: the left piece has $f \equiv 0$ so its mean is 0, and the right piece has mean $(1 - c)/(1 - s)$, giving $s(1 - s) \{(1 - c)/(1 - s)\}^2 = s(1 - c)^2/(1 - s)$, increasing in s on $[0, c]$. For $s \geq c$ the right piece has $f \equiv 1$ and the left has mean $(s - c)/s$, giving $s(1 - s) \{1 - (s - c)/s\}^2 = (1 - s)c^2/s$, decreasing in s . Both expressions equal $c(1 - c)$ at $s = c$. $\square$

Theorem 11.30 (A structure a tree can hold but greedy growing cannot find). Let $X = (X_1, X_2)$ be uniform on $[0, 1]^2$ and $f(\mathbf{x}) = 1$ when $x_1 > \frac{1}{2}$ and $x_2 > \frac{1}{2}$ both hold or both fail, and $f(\mathbf{x}) = 0$ otherwise, that is $f = 1$ on two opposite quarters and 0 on the other two. Then

- (i) a tree with four leaves reproduces f exactly;
- (ii) the best linear predictor of $f(X)$ is the constant $\frac{1}{2}$, with mean squared error $\frac{1}{4}$;
- (iii) every single axis-aligned split of $[0, 1]^2$ has population split gain exactly 0.

Proof. (i) Split at $x_1 = \frac{1}{2}$ and then each half at $x_2 = \frac{1}{2}$.

(ii) $\mathbf{E}[f] = \frac{1}{2}$. For the covariance with X_1 ,

$$\mathbf{E}[fX_1] = \int_{1/2}^1 \int_{1/2}^1 x_1 + \int_0^{1/2} \int_0^{1/2} x_1 = \frac{1}{2} \int_{1/2}^1 x \, dx + \frac{1}{2} \int_0^{1/2} x \, dx = \frac{1}{2} \cdot \frac{3}{8} + \frac{1}{2} \cdot \frac{1}{8} = \frac{1}{4},$$

which equals $\mathbf{E}[f] \mathbf{E}[X_1]$, so $\text{cov}(f, X_1) = 0$, and by symmetry $\text{cov}(f, X_2) = 0$. As in the proof of Theorem 11.29, the best $a + b_1X_1 + b_2X_2$ takes $a = \mathbf{E} f - b_1 \mathbf{E} X_1 - b_2 \mathbf{E} X_2$ and minimises $\text{var}(f - b_1X_1 - b_2X_2) = \text{var}(f) - 2 \sum_j b_j \text{cov}(f, X_j) + \sum_j b_j^2/12$, using that X_1, X_2 are independent with variance $1/12$; with both covariances 0 this is least at $b_1 = b_2 = 0$. The best linear predictor

is therefore the constant $\mathbf{E}[f] = \frac{1}{2}$ and its error is $\text{var}(f) = \frac{1}{4}$.

(iii) Take a split at $X_1 = s$ with $s \leq \frac{1}{2}$. On the left, $x_1 \leq s \leq \frac{1}{2}$, so $f = \mathbf{1}\{x_2 \leq \frac{1}{2}\}$ and the mean is $\frac{1}{2}$. On the right, conditionally on X_1 the mean of f is $\frac{1}{2}$ whichever side of $\frac{1}{2}$ X_1 falls, so the mean is $\frac{1}{2}$ there too. The two conditional means coincide and the gain, which by Unit 10 is a positive multiple of their squared difference, vanishes. The cases $s > \frac{1}{2}$ and splits on X_2 are the same computation. $\square$

Remark. Part (iii) is the limitation that matters and it is not the one usually stated. The target is a partition predictor with four boxes, exactly the class a tree fits, and the tree still fails: the greedy criterion of Unit 10 sees nothing at the first split, so whatever the algorithm does next is driven by sampling noise. What a tree can represent and what growing it finds are different questions.

3.3 Instability

Remark. A tree is a discontinuous function of the data. The split chosen at a node is the $\arg \max$ of finitely many gains, so an arbitrarily small perturbation that reverses the order of the top two changes the split, and every box below it. Nothing above bounds the resulting change in the fitted function.

Remark (Advantages and disadvantages; ISLR §8.1.4). ISLR lists for trees: they are easy to explain, even more so than linear regression, arguably closer to human decision-making, can be displayed graphically and read by non-experts, and handle qualitative predictors without dummy variables; against them, their predictive accuracy is generally below that of other methods, and they are not robust, a small change in the data changing the tree a great deal, which is the instability above.

4 Computation in R

ISLR §8.3.1–8.3.2 (labs). SRM learning outcomes 4(a), 4(b).

4.1 The ISLR labs

Example 11.1 (Classification trees on **Carseats**; ISLR §8.3.1). **Sales** of child car seats at 400 stores, recoded as **High = Yes** when **Sales** > 8; the full tree, then a tree grown on 200 training stores, pruned by cross-validation on the error count, with ISLR's seeds:

```
1 suppressMessages({library(ISLR2); library(tree); library(rpart)})
2 Car <- Carseats; Car$High <- factor(ifelse(Car$Sales <= 8, "No", "Yes"))
3 tc <- tree(High ~ . - Sales, Car); sc <- summary(tc)
4 c(leaves = sc$size, deviance = sc$dev, residual_df = sc$df, training_errors = sc$misclas
```

```
##           leaves      deviance    residual_df
##           27.0        170.7         373.0
## training_errors
##           36.0
```

```
1 set.seed(2); trn <- sample(1:nrow(Car), 200)
2 tt <- tree(High ~ . - Sales, Car, subset = trn)
3 acc <- function(m) mean(predict(m, Car[-trn, ], type = "class") == Car$High[-trn])
4 pc <- predict(tt, Car[-trn, ], type = "class"); table(predicted = pc, observed = Car$High
```

```
##           observed
```

```
## predicted  No Yes
##           No 104 33
##           Yes 13 50
```

```
1 set.seed(7); cm <- cv.tree(tt, FUN = prune.misclass); rbind(size = cm$size, cv_errors =
```

```
##           [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9]
## size        21  19  14   9   8   5   3   2   1
## cv_errors    75  75  75  74  82  83  83  85  82
```

```
1 pt <- prune.misclass(tt, best = 9)
2 c(full_accuracy = mean(pc == Car$High[-trn]), pruned9_accuracy = acc(pt),
3   pruned14_accuracy = acc(prune.misclass(tt, best = 14)))
```

```
##      full_accuracy  pruned9_accuracy pruned14_accuracy
##              0.770              0.775              0.770
```

ISLR's numbers are reproduced: 27 leaves, training error $36/400 = 9\%$ and residual mean deviance $170.66/373$; on the test stores $(104 + 50)/200 = 77\%$ correct; cross-validation chooses 9 leaves with 74 errors, and the pruned tree classifies 77.5% correctly, the 14-leaf tree 77%. The full-tree accuracy is computed from the same predictions as the table: `predict.tree` breaks a tie between equal class proportions in a leaf at random, so a second call can classify a store differently.

```

Exam set.seed(1); trb <- sample(1:nrow(Boston), nrow(Boston) / 2)
2  tbl <- tree(medv ~ ., Boston, subset = trb); sb <- summary(tbl)
3  as.character(sb$used); sb$size

```

```

## [1] "rm"      "lstat" "crim"  "age"
## [1] 7

```

```

1  set.seed(1); cvl <- cv.tree(tbl); rbind(size = cvl$size, cv_deviance = round(cvl$dev))

```

```

##           [,1] [,2] [,3] [,4] [,5] [,6] [,7]
## size           7    6    5    4    3    2    1
## cv_deviance 4337 4322 5070 5853 6561 9803 19697

```

```

1  yhb <- predict(tbl, newdata = Boston[-trb, ])
2  c(test_mse = mean((yhb - Boston$medv[-trb])^2),
3    root_mse = sqrt(mean((yhb - Boston$medv[-trb])^2)),
4    leaf_rm_above_7.553 = tbl$frame$yval[tbl$frame$var == "<leaf>"][7])

```

```

##           test_mse           root_mse leaf_rm_above_7.553
##           35.29           5.94           45.38

```

Four variables enter, **rm**, **lstat**, **crim** and **age**; tracts with **rm** ≥ 7.553 are predicted at \$45,400;

the test mean squared error of the seven-leaf tree is 35.29, a root mean squared error of about \$5,941. All are ISLR's. The cross-validated deviance is smallest at six leaves (4322), only 15 below the full tree's 4337; ISLR keeps the full tree, and the two are practically tied.

Example 11.3 (Pruning the `Hitters` tree; ISLR §8.1.1, Figures 8.4–8.5). `log(Salary)` on nine predictors, half the 263 players for training, the tree grown by `tree` and its cost-complexity sequence judged by six-fold cross-validation, by the training error and by the test error:

```

1 Hp <- na.omit(Hitters); Hp$lsal <- log(Hp$Salary)
2 set.seed(1); trh <- sample(nrow(Hp), 132)
3 th <- tree(lsal ~ Years + Hits + RBI + PutOuts + Walks + Runs + Assists + HmRun + Errors
4         Hp, subset = trh)
5 set.seed(1); cvh <- cv.tree(th, K = 6)
6 fitk <- function(k, d) if (k == 1) rep(mean(Hp$lsal[trh]), nrow(d)) else
7   predict(prune.tree(th, best = k), d)
8 round(rbind(size = cvh$size, cv_mse = cvh$dev / 132,
9   train = sapply(cvh$size, function(k) mean((fitk(k, Hp[trh, ]) - Hp$lsal[trh])^2)),
10  test = sapply(cvh$size, function(k) mean((fitk(k, Hp[-trh, ]) - Hp$lsal[-trh])^2))), 3

```

```
##           [,1]  [,2]  [,3]  [,4]  [,5]  [,6]  [,7]  [,8]  [,9]
```

```
## size    10.000  9.000  8.000  7.000  6.000  4.000  3.000  2.000  1.000
## cv_mse   0.447  0.438  0.418  0.420  0.446  0.433  0.445  0.471  0.792
## train    0.205  0.214  0.226  0.238  0.257  0.313  0.364  0.449  0.780
## test     0.367  0.375  0.387  0.389  0.379  0.361  0.347  0.429  0.796
```

The grown tree has 10 leaves. The training error falls with every leaf, as it must; the test error is smallest at 3 leaves, ISLR's choice; cross-validation is smallest at 8 leaves, though its curve is nearly flat from 3 to 10, where every value lies within 0.03 of the minimum. With 132 observations the cross-validation curve is noisy, and the size it picks is correspondingly uncertain.

Example 11.4 (A classification tree for **Heart**; ISLR §8.1.2, Figure 8.6). The 297 complete records of ISLR's **Heart** data, heart disease **AHD** on 13 predictors, among them the qualitative **Thal** and **ChestPain**:

```
1 Hd <- na.omit(read.csv("data/islr/Heart.csv", stringsAsFactors = TRUE)[, -1])
2 tHd <- tree(AHD ~ ., Hd); c(leaves = sum(tHd$frame$var == "<leaf>"))
```

```
## leaves
##      20
```

```
1 print(prune.misclass(tHd, best = 6), digits = 3)
```

```
## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
## 1) root 297 410.0 No ( 0.5387 0.4613 )
##    2) Thal: normal 164 175.0 No ( 0.7744 0.2256 )
##      4) Ca < 0.5 115  81.2 No ( 0.8870 0.1130 ) *
##      5) Ca > 0.5 49   67.9 No ( 0.5102 0.4898 )
##        10) ChestPain: nonanginal,nontypical,typical 29  32.1 No ( 0.7586 0.2414 ) *
##        11) ChestPain: asymptomatic 20  16.9 Yes ( 0.1500 0.8500 ) *
##    3) Thal: fixed,reversable 133 149.0 Yes ( 0.2481 0.7519 )
##      6) Ca < 0.5 59  81.4 Yes ( 0.4576 0.5424 )
##        12) ExAng < 0.5 33  42.0 No ( 0.6667 0.3333 ) *
##        13) ExAng > 0.5 26  25.5 Yes ( 0.1923 0.8077 ) *
##      7) Ca > 0.5 74  41.6 Yes ( 0.0811 0.9189 ) *
```

The unpruned tree has 20 leaves. Its first split is on **Thal**, a qualitative predictor split into two groups of its levels, as in ISLR's figure, and the six-leaf pruned tree uses only **Thal**, **Ca**, **ChestPain** and **ExAng**. ISLR points out splits in its tree whose two leaves predict the same class; such a split leaves the error count unchanged (Theorem 11.22) but lowers the Gini index or the deviance by making one leaf purer, which is why trees are grown on those and not on the error count.

The checks below use **Boston** in full ($n = 506$) for regression and **Car** for classification.

```
1 tb <- tree(medv ~ ., Boston)
2 c(n_boston = nrow(Boston), leaves_boston = sum(tb$frame$var == "<leaf>"))
```

```
##          n_boston leaves_boston
##              506              9
```

4.2 The pruning sequence (§1.2)

```
1 fr <- tb$frame; id <- as.integer(rownames(fr))
2 Rn <- setNames(fr$dev, id); lf <- setNames(fr$var == "<leaf>", id)
3 mval <- function(a, t) { k <- as.character(t)
4   if (lf[[k]]) Rn[[k]] + a else
5     min(Rn[[k]] + a, mval(a, 2*t) + mval(a, 2*t + 1)) }
6 alpha_t <- function(t) uniroot(function(a)
7   mval(a, 2*t) + mval(a, 2*t + 1) - Rn[[as.character(t)]] - a,
8   c(0, 1e7), tol = 1e-9)$root
9 manual <- sort(sapply(id[!lf], alpha_t))
10 pk <- prune.tree(tb)$k[-1]
11 signif(rbind(manual_thresholds = manual, package_k = pk,
```

```
12 difference = manual - pk), 5)
```

```
##           [,1]      [,2]      [,3]      [,4]
## manual_thresholds 4.455e+02 7.256e+02 1.137e+03 1.442e+03
## package_k        4.455e+02 7.256e+02 1.137e+03 1.442e+03
## difference        -1.137e-13 -3.411e-13 2.274e-13 -4.547e-13
##           [,5]      [,6]      [,7]      [,8]
## manual_thresholds 2.520e+03 3.061e+03 7.312e+03 19340
## package_k        2.520e+03 3.061e+03 7.312e+03 19340
## difference        9.095e-13 -9.095e-13 1.091e-11      0
```

```
1 leafset <- function(a, t) { k <- as.character(t)
2   if (lf[[k]] || Rn[[k]] + a <= mval(a, 2*t) + mval(a, 2*t + 1)) t
3   else c(leafset(a, 2*t), leafset(a, 2*t + 1)) }
4 grid <- c(0, (head(manual, -1) + tail(manual, -1))/2, max(manual) + 1)
5 t(sapply(grid, function(a)
6   c(alpha = a, manual_leaves = length(leafset(a, 1)),
7     package_leaves = sum(prune.tree(tb, k = a)$frame$var == "<leaf>"))))
```

```
##           alpha manual_leaves package_leaves
## [1,]      0.0              9              9
```

##	[2,]	585.5	8	8
##	[3,]	931.2	7	7
##	[4,]	1289.4	6	6
##	[5,]	1981.1	5	5
##	[6,]	2790.6	4	4
##	[7,]	5186.4	3	3
##	[8,]	13325.7	2	2
##	[9,]	19340.6	1	1

`rpart` reports the same thresholds divided by the cost of the root (Corollary 11.13).

```

1 rp <- rpart(medv ~ ., Boston, control = rpart.control(cp = 0, xval = 0))
2 R0 <- sum((Boston$medv - mean(Boston$medv))^2)
3 signif(rbind(cp = head(rp$cptable[, "CP"], 4),
4             cp_times_root_cost = head(rp$cptable[, "CP"], 4)*R0,
5             tree_thresholds = rev(manual)[1:4])), 5)

```

##		1	2	3	4
## cp		0.4527	0.1712	0.07166	0.03616
## cp_times_root_cost		19340.0000	7311.9000	3061.00000	1544.80000
## tree_thresholds		19340.0000	7311.9000	3061.00000	2520.30000

The first three agree exactly: both packages make the same first three splits, so they prune the same branches at the same prices. They part at the fourth because their stopping rules differ, **tree** using **mindev** and **minsize** where **rpart** uses **minsplit** and **minbucket**, so the trees being pruned are not the same tree. Figure 11.1 draws the lower envelope and the leaf count of the **Boston** pruning sequence.




Figure 11.1: Left: for each subtree in the sequence, the line $\alpha \mapsto R(S) + \alpha|S|$, with the value function $m_\alpha(r)$ of Definition 11.5 as their lower envelope (Lemma 11.7), concave and piecewise linear; the dotted verticals are the thresholds of Theorem 11.8. Right: the leaf count of $T(\alpha)$, a decreasing step function (Theorem 11.11) jumping here at every threshold (in general at some, Corollary 11.12).

4.3 Cross-validation (§1.3)

`cv.tree` forms its folds by drawing a fold label for each observation independently, so the folds are not of equal size.

```
1 set.seed(3); K <- 10; n <- nrow(Boston)
2 fold <- sample(K, n, replace = TRUE)
3 c(intended_size = n/K, smallest = min(table(fold)),
4   largest = max(table(fold)))
```

```
## intended_size      smallest      largest
##           50.6           38.0           60.0
```

Each fold's tree is pruned at the **full-data** thresholds, so the subtree a fold contributes at a given α need not have the number of leaves printed against it.

```
1 mf <- model.frame(tb)
2 sz <- sapply(1:4, function(i) {
3   t1 <- tree(model = mf[fold != i, , drop = FALSE])
4   sapply(manual, function(a) sum(prune.tree(t1, k = a)$frame$var == "<leaf>")) })
5 rbind(alpha = round(manual), full_data_size = prune.tree(tb)$size[-1], t(sz))
```

```
##           [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8]
```

## alpha	445	726	1137	1442	2520	3061	7312	19340
## full_data_size	8	7	6	5	4	3	2	1
##	8	7	5	4	4	4	2	1
##	7	7	6	5	4	3	2	1
##	10	8	6	5	5	3	2	1
##	9	7	6	6	5	3	3	1

```
1 set.seed(3); cvb <- cv.tree(tb, K = 10)
2 rbind(size = cvb$size, cv_deviance = round(cvb$dev))
```

##	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]
## size	9	8	7	6	5	4	3	2
## cv_deviance	12767	13071	13350	13442	15607	17788	18590	27193
##	[,9]							
## size	1							
## cv_deviance	42844							

Cross-validation keeps the whole tree here: the deviance is smallest at nine leaves, the largest the growing rule produced. Pruning is not always an improvement, and the sequence is a family to search, not a prescription to shrink.

The one standard error rule of Definition 11.15 needs a standard error, which **rpart** supplies as **xstd**.

```

1 set.seed(5)
2 rp2 <- rpart(medv ~ ., Boston, control = rpart.control(cp = 0.001, xval = 10))
3 ct <- rp2$cptable; j <- which.min(ct[, "xerror"])
4 thr <- ct[j, "xerror"] + ct[j, "xstd"]
5 c(minimising_nsplitt = ct[j, "nsplitt"], min_xerror = ct[j, "xerror"],
6   one_se_threshold = thr,
7   one_se_nsplitt = ct[which(ct[, "xerror"] <= thr)[1], "nsplitt"])

```

```

## minimising_nsplitt      min_xerror  one_se_threshold
##           23.0000          0.2443          0.2804
##      one_se_nsplitt
##           7.0000

```

4.4 Impurity (§2.2)

```

1 Mf <- function(p) 1 - max(p); Gf <- function(p) 1 - sum(p^2)
2 Hf <- function(p) { p <- p[p > 0]; -sum(p*log(p)) }
3 set.seed(1); P <- matrix(rexp(4000*4), ncol = 4); P <- P/rowSums(P)
4 c(smallest_G_minus_M = min(apply(P, 1, Gf) - apply(P, 1, Mf)),
5   smallest_H_minus_G = min(apply(P, 1, Hf) - apply(P, 1, Gf)),
6   largest_G_minus_2M_two_classes = max(sapply(seq(0, 1, 0.001),

```

```

7  function(p) Gf(c(p, 1 - p)) - 2*Mf(c(p, 1 - p))))

##          smallest_G_minus_M          smallest_H_minus_G
##          0.01403                0.07013
## largest_G_minus_2M_two_classes
##          0.00000

```

Two splits of a node holding 400 of each class: A into (300, 100) and (100, 300), B into (200, 400) and (200, 0).

```

1  wt <- function(nL, nR, f) { N <- sum(nL) + sum(nR)
2    sum(nL)/N*f(nL/sum(nL)) + sum(nR)/N*f(nR/sum(nR)) }
3  imp <- function(nL, nR) c(M = wt(nL, nR, Mf), G = wt(nL, nR, Gf),
4    H = wt(nL, nR, Hf))
5  round(rbind(parent = c(M = Mf(c(.5, .5)), G = Gf(c(.5, .5)),
6    H = Hf(c(.5, .5))),
7    split_A = imp(c(300, 100), c(100, 300)),
8    split_B = imp(c(200, 400), c(200, 0))), 4)

```

```

##          M          G          H
## parent  0.50 0.5000 0.6931
## split_A 0.25 0.3750 0.5623

```

```
| ## split_B 0.25 0.3333 0.4774
```

Both splits reduce M from 0.50 to 0.25, so the error count cannot choose between them; G and H both prefer B, which produces a pure leaf. Proposition 11.25 identifies what **tree** prints for a classification tree.

```
1 lc <- tc$frame$var == "<leaf>"
2 c(package_deviance = deviance(tc),
3   manual_2_sum_n_H = 2*sum(tc$frame$n[lc]*apply(tc$frame$yprob[lc, ], 1, Hf)),
4   training_errors = sum(tc$frame$n[lc]*apply(tc$frame$yprob[lc, ], 1, Mf)))
```

```
| ## package_deviance manual_2_sum_n_H training_errors
| ##                170.7             170.7          36.0
```




Figure 11.2: The three impurities for two classes. Theorem 11.23 places M below G below H and traps G between M and $2M$; all three vanish at $p = 0$ and $p = 1$ and peak at $p = \frac{1}{2}$. Only M has a corner there, and it is affine on each side, which is Theorem 11.22 in one picture.

4.5 Categorical predictors (§2.4)

```

1 set.seed(4); m <- 300; L <- 8
2 lv <- factor(sample(letters[1:L], m, TRUE))
3 yc <- c(2, 9, 4, 7, 1, 5, 8, 3)[as.integer(lv)] + rnorm(m, 0, 2)
4 Q <- function(A) { i <- lv %in% levels(lv)[A]
5   if (!any(i) || all(i)) Inf else
6     sum((yc[i] - mean(yc[i]))^2) + sum((yc[!i] - mean(yc[!i]))^2) }
7 subs <- lapply(1:(2^(L - 1) - 1),
8   function(b) which(as.integer(intToBits(b))[1:L] == 1))
9 ord <- order(tapply(yc, lv, mean))
10 c(subsets_tried = length(subs), ordered_tried = L - 1,
11   best_exhaustive = min(sapply(subs, Q)),
12   best_ordered = min(sapply(1:(L - 1), function(k) Q(ord[1:k]))),
13   difference = min(sapply(1:(L - 1), function(k) Q(ord[1:k])) -
14     min(sapply(subs, Q)))

```

```

## subsets_tried ordered_tried best_exhaustive
##           127             7           1668
## best_ordered difference
##           1668             0

```

4.6 Trees against linear models (§3.1–§3.2)

With $\beta_1 = 3$ the predicted error is $9/(12J^2)$.

```
1 set.seed(8); N <- 2e5; xu <- runif(N)
2 t(sapply(c(2, 4, 8, 16), function(J) {
3   b <- cut(xu, seq(0, 1, length.out = J + 1))
4   c(J = J, mse = mean((3*xu - ave(3*xu, b))^2), theory = 9/(12*J^2)) })))
```

```
##           J      mse  theory
## [1,]    2 0.187635 0.18750
## [2,]    4 0.046878 0.04688
## [3,]    8 0.011719 0.01172
## [4,]   16 0.002925 0.00293
```

```
1 fs <- as.numeric(xu > 0.5)
2 c(linear_mse = mean(residuals(lm(fs ~ xu))^2),
3   theory = 0.25*(1 - 3*0.25),
4   two_leaf_tree_mse = mean((fs - ave(fs, xu > 0.5))^2))
```

```
##           linear_mse           theory two_leaf_tree_mse
##           0.0625           0.0625           0.0000
```

On a sample of 400 the first-split gains are noise.

```

1 set.seed(11); q <- 400; z1 <- runif(q); z2 <- runif(q)
2 fx <- as.numeric((z1 > .5) == (z2 > .5)); yx <- fx + rnorm(q, 0, .3)
3 gn <- function(s, v) { i <- if (v == 1) z1 < s else z2 < s
4   sum(i)*sum(!i)/q*(mean(yx[i]) - mean(yx[!i]))^2 }
5 ss <- c(.2, .35, .5, .65, .8)
6 round(rbind(cut = ss, gain_x1 = sapply(ss, gn, v = 1),
7   gain_x2 = sapply(ss, gn, v = 2),
8   percent_of_total_x2 = 100*sapply(ss, gn, v = 2)/
9     sum((yx - mean(yx))^2)), 4)

```

##	[,1]	[,2]	[,3]	[,4]	[,5]
## cut	0.2000	0.3500	0.5000	0.6500	0.8000
## gain_x1	0.7533	0.2945	0.0012	0.5719	0.8826
## gain_x2	2.2278	0.9565	0.5922	0.6172	0.3380
## percent_of_total_x2	1.7156	0.7366	0.4560	0.4753	0.2603

The largest gain in the table is at $x_2 < 0.2$ and is 1.7% of the total sum of squares, while splitting x_1 at the cut that matters, $\frac{1}{2}$, has the smallest gain of the ten. A greedy tree splits where the noise is. Figure 11.3 draws the staircase error and the checkerboard.




Figure 11.3: Left: Theorem 11.28, the error of the J -box grid approximation to a linear truth, on a log scale; the decay $J^{-2/p}$ flattens as the dimension grows. Right: the checkerboard truth of Theorem 11.30, which four boxes fit exactly and which no single split improves at all.

4.7 Instability (§3.3)

```

1 set.seed(9); Rt <- replicate(200, {
2   i <- sample(nrow(Boston), nrow(Boston)/2)
3   a <- tree(medv ~ ., Boston[i, ]); b <- tree(medv ~ ., Boston[-i, ])
4   c(as.character(a$frame$var[1]), as.character(b$frame$var[1]),
5     sqrt(mean((predict(a, Boston) - predict(b, Boston))^2))) })
6 sort(table(c(Rt[1, ], Rt[2, ])), decreasing = TRUE)

```

```

##
##      rm lstat
## 215    185

```

```

1 c(agree_on_root_variable = mean(Rt[1, ] == Rt[2, ]),
2   rms_prediction_difference = mean(as.numeric(Rt[3, ])),
3   sd_of_medv = sd(Boston$medv))

```

```

##      agree_on_root_variable rms_prediction_difference
##                0.185                4.124
##                sd_of_medv
##                9.197

```

The two halves choose `rm` and `lstat` in nearly equal numbers, agree on the root variable under a fifth of the time, and their predictions on the same observations differ in root mean square by 45% of the standard deviation of the response itself.

4.8 Growing, pruning and selecting in practice

Stopping early instead of pruning

Growing while the gain exceeds a threshold and pruning a large tree are not the same procedure. The checkerboard truth of Theorem 11.30 is the extreme case: no first split has any population gain at all, so every threshold above zero stops at the root.

```
1 set.seed(15); q <- 800; w1 <- runif(q); w2 <- runif(q)
2 fw <- as.numeric((w1 > .5) == (w2 > .5))
3 dw <- data.frame(w1, w2, y = fw + rnorm(q, 0, .3))
4 t(sapply(c(0.05, 0.01, 0.002, 0), function(md) {
5   g <- tree(y ~ w1 + w2, dw,
6             control = tree.control(q, mindev = md, minsize = 10))
7   c(mindev = md, grown_leaves = sum(g$frame$var == "<leaf>"),
8     mse_of_fit = mean((predict(g, dw) - fw)^2)) })))
```

```
##      mindev grown_leaves mse_of_fit
## [1,]  0.050           1  0.250011
## [2,]  0.010           1  0.250011
## [3,]  0.002          11  0.004584
## [4,]  0.000         130  0.037224
```

```
1 c(variance_of_f = mean((fw - mean(fw))^2))
```

```
## variance_of_f
##           0.25
```

At the default `mindev` of 0.01 the tree is the root and its mean squared error is the variance of the target: nothing has been learnt. Growing without a threshold and pruning afterwards does better than either extreme.

```
1 g0 <- tree(y ~ w1 + w2, dw, control = tree.control(q, mindev = 0, minsize = 10))
2 t(sapply(c(2, 4, 8, 16, 40, 130), function(J) {
3   pj <- prune.tree(g0, best = J)
4   c(requested = J, leaves = sum(pj$frame$var == "<leaf>"),
5     mse = mean((predict(pj, dw) - fw)^2)) }))
```

```
##      requested leaves      mse
```

## [1,]	2	5	0.052843
## [2,]	4	5	0.052843
## [3,]	8	8	0.002027
## [4,]	16	18	0.008559
## [5,]	40	48	0.021706
## [6,]	130	130	0.037224

The symptom is a tree with one leaf on data whose structure is known to be richer, and a fit that improves sharply the moment the stopping rule is loosened. Grow past the point where splits look worthwhile, then prune (Algorithm 8.1 of ISLR, §1). The best pruned tree here has mean squared error 0.0020 against the 0.2500 the default stopping rule delivers, a factor of 123, and against 0.0372 for the fully grown tree that was pruned. Both halves are needed: growing finds the structure, pruning removes what the growing added afterwards. Note also that **best** requests a size and returns the smallest tree in the sequence with at least that many leaves, so 2 and 4 both return five leaves.

Pruning with the criterion used for growing

Both the error count and the deviance are legitimate node costs (Proposition 11.3), and they select different trees.

```

1 tf <- tree(High ~ . - Sales, Car)
2 t(sapply(1:6, function(s) { set.seed(s)
3   a <- cv.tree(tf, FUN = prune.misclass); b <- cv.tree(tf)
4   c(seed = s, chosen_by_error_count = a$size[which.min(a$dev)],
5     chosen_by_deviance = b$size[which.min(b$dev)]) })))

```

##	seed	chosen_by_error_count	chosen_by_deviance
## [1,]	1	17	3
## [2,]	2	27	2
## [3,]	3	12	3
## [4,]	4	7	3
## [5,]	5	27	3
## [6,]	6	27	2

The symptom is two runs of the same pipeline reporting trees of two and of twenty-seven leaves. Prune on the loss the model will be judged by. The deviance of Proposition 11.25 is a log likelihood and punishes a confident wrong leaf without limit; the error count of Proposition 11.17 does not. If the tree will be scored by its error rate, use **FUN = prune.misclass**. Note also that the error-count choice is itself unstable across folds, ranging over most of the sequence as the seed changes, which is

the instability of §3.3 appearing in the selection step.

Reading the training curve

```
1 pr <- prune.tree(tb)
2 rbind(leaves = pr$size, training_deviance = round(pr$dev),
3       cv_deviance = round(cvb$dev))
```

##	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]
## leaves	9	8	7	6	5	4	3
## training_deviance	6734	7179	7905	9042	10484	13004	16065
## cv_deviance	12767	13071	13350	13442	15607	17788	18590
##	[,8]	[,9]					
## leaves	2	1					
## training_deviance	23377	42716					
## cv_deviance	27193	42844					

The symptom is a cost that falls monotonically with the number of leaves, used to choose the number of leaves. The training cost is decreasing in the tree by the refinement property (Proposition 11.3), so its minimiser is always $T_{\max}$ and it carries no information about size. Only the cross-validated column

can be minimised. On these data both columns happen to point at nine leaves, which is a coincidence of this example and not a reason to trust the first.

Cross-validating the pruning but not the growing

```
1 set.seed(21); B <- 40; n <- nrow(Boston)
2 res <- t(replicate(B, {
3   f <- sample(10, n, replace = TRUE)
4   full <- tree(medv ~ ., Boston)
5   honest <- leaky <- 0
6   for (k in 1:10) {
7     ho <- Boston[f == k, ]
8     h <- prune.tree(tree(medv ~ ., Boston[f != k, ]), best = 5)
9     honest <- honest + sum((ho$medv - predict(h, ho))^2)
10    leaky <- leaky + sum((ho$medv - predict(prune.tree(full, best = 5), ho))^2)
11  }
12  c(honest, leaky)/n })))
13 c(honest_cv_mse = mean(res[, 1]), leaky_cv_mse = mean(res[, 2]),
14    ratio = mean(res[, 2])/mean(res[, 1]))
```

```
## honest_cv_mse   leaky_cv_mse          ratio
##           29.6358       20.7186       0.6991
```

The symptom is a cross-validation loop in which only the pruning, or only the choice of α , is redone inside the fold while the splits come from a tree fitted to everything. The leaky loop above is the extreme case, in which nothing is refit inside the fold: it reports the training error of the full-data five-leaf tree. Refit from the raw data inside each fold, as Definition 11.14 requires. The splits are chosen by looking at the response, so a tree grown on all the data has already seen every held-out observation. The optimism is not small here: the leaky loop reports 30% less error than the honest one, which would turn a mediocre tree into a good one on paper.

Reading a single tree as a statement about variables

```
1 set.seed(31); imp <- replicate(200, {
2   i <- sample(nrow(Boston), replace = TRUE)
3   as.character(tree(medv ~ ., Boston[i, ])$frame$var[1]) })
4 round(sort(prop.table(table(imp))), decreasing = TRUE), 3)
```

```
## imp
```

```
##      rm lstat  
## 0.575 0.425
```

The symptom is a report that names the root split variable as the most important predictor. Resample and look at the distribution. Here two variables share the root almost equally, and which one appears is a property of the sample and not of the population. Importance measures averaged over many trees are the only form in which the question has a stable answer. Figure [11.4](#) shows how much of the five-leaf **Boston** tree the first split carries.




Figure 11.4: Left, **Boston** pruned to five leaves with branch lengths proportional to the drop in the node cost, so a long branch is a split that mattered: the first split on **rm** accounts for most of the tree. Right, the **Carseats** tree of Example 11.1 drawn with uniform branches, which nine leaves need to stay legible. It splits on **ShelveLoc** directly, with no dummy coding.

5 Exercises

Exercise 11.1 (ISLR §8.4, exercise 3). With two classes, plot the Gini index, the classification error and the entropy as functions of $\hat{p}_{m1}$.

Solution. With $\hat{p}_{m1} = p$, $M = \min\{p, 1-p\}$, $G = 2p(1-p)$ and $H = -p \ln p - (1-p) \ln(1-p)$ (Definition 11.19); Figure 11.2 is the plot. All three vanish at $p \in \{0, 1\}$ and peak at $p = \frac{1}{2}$ with values $\frac{1}{2}$, $\frac{1}{2}$ and $\ln 2 = 0.693$, and $M \leq G \leq H$ throughout (Theorem 11.23); only M has a corner at $\frac{1}{2}$, the source of Theorem 11.22. ■

Exercise 11.2 (ISLR §8.4, exercise 6). Explain in detail the algorithm used to fit a regression tree.

Solution. (1) Grow: starting from the whole predictor space, choose the predictor j and cutpoint s minimising the residual sum of squares of the two resulting boxes (Unit 10's split criterion, equivalently maximising $n_L n_R (\bar{y}_L - \bar{y}_R)^2 / n$), and repeat within each box until each contains fewer than a minimum number of observations; each leaf predicts its mean. (2) Prune: for each $\alpha \geq 0$ the subtree minimising $R(S) + \alpha |S|$ (Definition 11.4) is the canonical $T(\alpha)$ of Theorem 11.10,

and as α grows these form a finite nested sequence (Corollary 11.12) computed by the recursion of Lemma 11.6. (3) Choose α by cross-validation over folds, regrowing and pruning in each fold (Definition 11.14). (4) Return $T(\hat{\alpha})$ from the tree grown on all the data. ■

Exercise 11.3 (ISLR §8.4, exercise 8, parts (a)–(c)). For **Carseats** with **Sales** as a quantitative response: (a) split into training and test sets; (b) fit a regression tree, interpret it and give the test mean squared error; (c) choose the tree size by cross-validation and say whether pruning improves the test mean squared error. Parts (d)–(f) are bagging, random forests and BART, Unit 12's and beyond the syllabus.

Solution.

```
1 set.seed(1); tr8 <- sample(nrow(Carseats), 200)
2 ts8 <- tree(Sales ~ ., Carseats, subset = tr8)
3 mse8 <- function(m) mean((predict(m, Carseats[-tr8, ]) - Carseats$Sales[-tr8])^2)
4 as.character(summary(ts8)$used); c(leaves = summary(ts8)$size, test_mse = mse8(ts8))
```

```
## [1] "ShelveLoc"    "Price"        "Age"          "Advertising"
## [5] "CompPrice"    "US"
##  leaves test_mse
```

```
##      18.000      4.922
```

```
1 set.seed(1); cv8 <- cv.tree(ts8); rbind(size = cv8$size, cv_deviance = round(cv8$dev, 1))
```

```
##           [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9]
## size      18.0  17   16   15   14   13   12   11   10
## cv_deviance 984.4 1031 1036 1027 1027 1056 1045 1061 1061
##           [,10] [,11] [,12] [,13] [,14] [,15] [,16] [,17]
## size           8     7     6     5     4     3     2     1
## cv_deviance  1226  1221  1219  1232  1337  1300  1338  1605
```

```
1 round(sapply(c(18, 8, 5), function(k) mse8(prune.tree(ts8, best = k))), 3)
```

```
## [1] 4.922 5.113 5.186
```

(a) 200 stores train and 200 test. (b) The 18-leaf tree splits first on **ShelveLoc** and then on **Price**, **Age**, **Advertising**, **CompPrice** and **US**: good shelving and low prices raise sales. Its test mean squared error is 4.92. (c) Cross-validation selects the full 18 leaves; pruning to 8 or 5 leaves raises the test error to 5.11 and 5.19, so pruning does not help here. ■

Exercise 11.4 (ISLR §8.4, exercise 9). For `OJ`, with `Purchase` the response: (a) take 800 training observations; (b) fit a tree and give the training error rate and number of leaves; (c) interpret one leaf of the printed tree; (d) plot and interpret it; (e) give the test confusion matrix and error rate; (f)–(h) choose the size by cross-validation on the error count; (i) prune to that size, or to five leaves if the full tree is chosen; (j), (k) compare training and test error rates of the pruned and unpruned trees.

Solution.

```
1 set.seed(1); tro <- sample(nrow(OJ), 800); to <- tree(Purchase ~ ., OJ, subset = tro)
2 so <- summary(to); c(leaves = so$size, training_errors = so$misclass[1],
3   training_error_rate = so$misclass[1] / 800)
```

```
##           leaves      training_errors training_error_rate
##           9.0000           127.0000           0.1588
```

```
1 print(to, digits = 4)
```

```
## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
```

```
## 1) root 800 1073.00 CH ( 0.60625 0.39375 )
## 2) LoyalCH < 0.5036 365 441.60 MM ( 0.29315 0.70685 )
## 4) LoyalCH < 0.280875 177 140.50 MM ( 0.13559 0.86441 )
## 8) LoyalCH < 0.0356415 59 10.14 MM ( 0.01695 0.98305 ) *
## 9) LoyalCH > 0.0356415 118 116.40 MM ( 0.19492 0.80508 ) *
## 5) LoyalCH > 0.280875 188 258.00 MM ( 0.44149 0.55851 )
## 10) PriceDiff < 0.05 79 84.79 MM ( 0.22785 0.77215 )
## 20) SpecialCH < 0.5 64 51.98 MM ( 0.14062 0.85938 ) *
## 21) SpecialCH > 0.5 15 20.19 CH ( 0.60000 0.40000 ) *
## 11) PriceDiff > 0.05 109 147.00 CH ( 0.59633 0.40367 ) *
## 3) LoyalCH > 0.5036 435 337.90 CH ( 0.86897 0.13103 )
## 6) LoyalCH < 0.764572 174 201.00 CH ( 0.73563 0.26437 )
## 12) ListPriceDiff < 0.235 72 99.81 MM ( 0.50000 0.50000 )
## 24) PctDiscMM < 0.196196 55 73.14 CH ( 0.61818 0.38182 ) *
## 25) PctDiscMM > 0.196196 17 12.32 MM ( 0.11765 0.88235 ) *
## 13) ListPriceDiff > 0.235 102 65.43 CH ( 0.90196 0.09804 ) *
## 7) LoyalCH > 0.764572 261 91.20 CH ( 0.95785 0.04215 ) *
```

```
1 po <- predict(to, OJ[-tro, ], type = "class"); table(predicted = po, observed = OJ$Purch
```

```
## observed
```

```
## predicted  CH  MM
##           CH 160 38
##           MM   8 64
```

```
1 set.seed(1); cvo <- cv.tree(to, FUN = prune.misclass); rbind(size = cvo$size, cv_errors = cvo$cv.errors)
```

```
##           [,1] [,2] [,3] [,4] [,5] [,6]
## size          9   8   7   4   2   1
## cv_errors    145  145  146  146  167  315
```

```
1 pp <- prune.misclass(to, best = 5); sp <- summary(pp)
2 round(c(pruned_leaves = sp$size, unpruned_train = so$misclass[1] / 800,
3         pruned_train = sp$misclass[1] / 800, unpruned_test = mean(po != OJ$Purchase[-tro]),
4         pruned_test = mean(predict(pp, OJ[-tro, ], type = "class") != OJ$Purchase[-tro]))
```

```
## pruned_leaves unpruned_train pruned_train unpruned_test
##          7.0000          0.1588          0.1625          0.1704
## pruned_test
##          0.1630
```

(b) Nine leaves, training error $127/800 = 15.9\%$; only LoyalCH, PriceDiff, SpecialCH, ListPriceDiff and PctDiscMM are used. (c) Leaf 8, LoyalCH < 0.0356 : 59 customers, 98.3%

of whom bought Minute Maid, so the leaf predicts **MM** with near certainty and deviance 10.14. (d) Brand loyalty **LoyalCH** makes the first three splits; prices matter only for customers of middling loyalty. (e) The test error rate is $46/270 = 17.0\%$. (f)–(h) The cross-validated error count is smallest, 145, at 9 and 8 leaves; the full tree is among the minimisers. (i) The request for five leaves returns the smallest size of at least five in the pruning sequence, seven (Corollary 11.12). (j) The pruned tree's training error is higher, 16.3% against 15.9%, as Proposition 11.3 requires. (k) Its test error is lower, 16.3% against 17.0%. ■

Exercise 11.5 (SOA sample question 9). Of 100 policies, 30 lapsed. Split 1 gives nodes of 20 (with 12 lapses) and 80 (with 18); split 2 gives nodes of 10 (with 8) and 90 (with 22). Totals are node impurities weighted by node size. Which is true? I. Split 1 is preferred on the total Gini index. II. Split 1 is preferred on the total entropy. III. Split 1 is preferred on the number of classification errors. (A) I only; (B) II only; (C) III only; (D) I, II and III; (E) none of these.

Solution. By Definition 11.19:

```
1 Gs <- function(n1, n) { p <- n1 / n; n * 2 * p * (1 - p) }
2 Hs <- function(n1, n) { p <- c(n1, n - n1) / n; -n * sum(ifelse(p > 0, p * log(p), 0)) }
```

```
3 Ms <- function(n1, n) min(n1, n - n1)
4 sp <- list(split1 = c(12, 20, 18, 80), split2 = c(8, 10, 22, 90))
5 round(t(sapply(sp, function(v) c(Gini = (Gs(v[1], v[2]) + Gs(v[3], v[4])) / 100,
6   entropy = (Hs(v[1], v[2]) + Hs(v[3], v[4])) / 100, errors = Ms(v[1], v[2]) + Ms(v[3],
```

```
##           Gini entropy errors
## split1 0.3750  0.5611      26
## split2 0.3644  0.5506      24
```

Split 2 has the smaller total Gini index (0.3644 against 0.375), entropy (0.5506 against 0.5611) and error count (24 against 26), so I, II and III are all false: (E). ■

Exercise 11.6 (SOA sample question 25). Which is true of tree pruning? I. Recursive binary splitting can overfit the data. II. A tree with more splits tends to have lower variance. III. With cost complexity pruning, $\alpha = 0$ gives a very large tree. (A) none; (B) I and II only; (C) I and III only; (D) II and III only; (E) none of these.

Solution. I: splitting reduces the training cost at every step (Proposition 11.3), so growing to small leaves fits noise. II is false: more splits mean leaves estimated from fewer observations, hence

more variance. III: at $\alpha = 0$ the minimum of R is attained by $T_{\max}$ (proof of Theorem 11.8), so the selected tree is as large as the growing produced, up to branches that reduce no cost (Corollary 11.12). (C). ■

Exercise 11.7 (SOA sample question 26). In each of the three panels the plane is shaded in two classes. Which can a classification tree fit with no error?

I



II



III




(A) I only; (B) II only; (C) III only; (D) I, II and III; (E) none of these.

Solution. A tree's regions are boxes cut by axis-parallel lines (Unit 10's definition of a split). Panel I is three boxes cut by two axis-parallel lines; II has a curved boundary and III a diagonal one, which finitely many axis-parallel cuts cannot reproduce exactly. (A). ■

Exercise 11.8 (SOA sample question 29). Which considerations may make trees preferable? I. Trees are easily interpretable. II. Trees can be displayed graphically. III. Trees are easier to explain than linear regression. (A) none; (B) I and II only; (C) I and III only; (D) II and III only; (E) none of these.

Solution. All three are among ISLR's advantages of trees (the remark of §3.3), so (A)–(D) are all wrong: (E). ■

Exercise 11.9 (SOA sample question 33). The regression tree below predicts the log claim amount (LCA) from age category (`agecat`) and vehicle age (`veh_age`).





Rank the predictions for I (**agecat** 1, **veh_age** 4), II (**agecat** 5, **veh_age** 5) and III (**agecat** 5, **veh_age** 3): (A) I < II < III; (B) I < III < II; (C) II < I < III; (D) II < III < I; (E) III < II <

I.

Solution. I fails $\text{agecat} \geq 1.5$ and lands in the right leaf, 8.146. II and III satisfy it and have $\text{veh_age} \geq 2.5$ and $\text{agecat} \geq 4.5$; then $\text{veh_age} < 3.5$ separates III (7.771) from II (8.028). So $\text{III} < \text{II} < \text{I}$: (E). ■


Exercise 11.10 (SOA sample question 48). Which partition represents the tree below (the left branch is the true inequality)?





Solution. The root $X_1 < t_1$ is a horizontal line at height t_1 (X_1 is on the vertical axis); the second split $X_2 < t_2$ acts only below it, a vertical segment up to t_1 ; the third, $X_1 < t_3$, acts only to the left of that segment, and its true branch is the leaf C. Only (B) has all three in that nesting with C below t_3 . ■

Exercise 11.11 (SOA sample question 51). The tree below predicts the weight of ducks.




Order the predictions for X (wing span 5.5 cm, male, age 7), Y (5.8 cm, female, age 5) and Z (5.7 cm, male, age 8): (A) $X < Y < Z$; (B) $X < Z < Y$; (C) $Y < X < Z$; (D) $Y < Z < X$; (E) $Z < X < Y$.

Solution. X: age ≤ 7.5 , male: 0.90 kg. Y: age ≤ 7.5 , female: 0.80 kg. Z: age > 7.5 , male, wing span > 5.3 : 1.25 kg. So $Y < X < Z$: (C). ■

Exercise 11.12 (SOA sample question 57). For the nine observations below and the regression tree shown, calculate the mean response of each leaf.

R	4.75	4.67	4.67	4.56	4.53	3.91	3.90	3.90	3.89
X	M	F	M	F	M	F	F	M	M
Y	A	A	D	D	B	C	B	D	B
Z	2	4	1	3	2	2	5	5	1



(A) 4.39, 4.38, 4.29, 3.90; (B) 4.26, 4.38, 4.62, 3.90; (C) 4.26, 4.39, 3.90, 4.29; (D) 4.64, 4.29, 4.38, 3.90; (E) 4.64, 4.38, 4.39, 3.90 (for T1, T2, T3, T4).

Solution. By Unit 10's proposition **the region constants are the region means**, each leaf predicts the mean of its observations:

```

1 R <- c(4.75, 4.67, 4.67, 4.56, 4.53, 3.91, 3.90, 3.90, 3.89)
2 X <- c("M", "F", "M", "F", "M", "F", "F", "M", "M"); Y <- c("A", "A", "D", "D", "B", "C", "B", "D", "B")
3 Z <- c(2, 4, 1, 3, 2, 2, 5, 5, 1)
4 leaf <- ifelse(Z <= 3, ifelse(Y %in% c("A", "B"), "T1", "T2"), ifelse(X == "F", "T3", "T4"))


```

```
5 round(tapply(R, leaf, mean), 2)
```

```
##      T1      T2      T3      T4  
## 4.39 4.38 4.28 3.90
```

T1 holds observations 1, 5, 9; T2 holds 3, 4, 6; T3 holds 2, 7, with mean $(4.67 + 3.90)/2 = 4.285$, which R prints as 4.28 and the answer key rounds to 4.29; T4 holds 8: (A). ■


Exercise 11.13 (SOA sample question 63). For the regression tree below of car-seat sales (thousands), predict sales for a store with **ShelveLoc** Good, **Price** 120, **Age** 57, **Advertising** 12.



(A) 4.7; (B) 6.6; (C) 7.4; (D) 9.2; (E) 9.4.

Solution. `ShelveLoc` is not Bad or Medium, so the right branch; `Price` = 120 $\geq$ 110, so its left leaf: 9.2 thousand. (D). ■

Exercise 11.14 (SOA sample question 66). Which set of regions represents the regression tree below (left branches are the true inequalities)?




(A) $R_3 = \{X_1 > 5\}$; (B) $R_3 = \{X_1 > 5, X_2 \leq 60\}$; (C) $R_1 = \{X_1 \leq 5, X_2 > 40\}$, $R_2 = \{X_1 \leq 5, X_2 \leq 40\}$; (D) $R_1 = \{X_1 \leq 5, X_2 \leq 40\}$, $R_2 = \{X_1 \leq 5, X_2 > 40\}$, $R_3 = \{5 < X_1 \leq 10\}$,

$R_4 = \{X_1 > 10, X_2 \leq 60\}$, $R_5 = \{X_1 > 10, X_2 > 60\}$; (E) as (B) with R_4 and R_5 exchanged.

Solution. Each region is the intersection of the inequalities on its path from the root: R_1 and R_2 lie under $X_1 \leq 5$ and divide at $X_2 = 40$; R_3 is $X_1 > 5$ and $X_1 \leq 10$; R_4 and R_5 lie under $X_1 > 10$ and divide at $X_2 = 60$. That is (D); (A), (B) and (E) give R_3 the wrong box and (C) exchanges R_1 and R_2 . ■

Exercise 11.15 (SOA sample question 73). The plot shows cross-validation, test and training mean squared errors against tree size under cost complexity pruning. Determine the size of the optimal pruned tree.



(A) 2; (B) 5; (C) 6; (D) 9; (E) 10.

Solution. The size is chosen by minimising the cross-validated error (Definition 11.14), not the training error, which decreases in the size (Proposition 11.3), nor the test error, which is not available when choosing. The cross-validation curve is lowest at 5: (B). ■